Self-Driving Car Engineer Nanodegree

Deep Learning

Project: Build a Traffic Sign Recognition Classifier

In this notebook, a template is provided for you to implement your functionality in stages, which is required to successfully complete this project. If additional code is required that cannot be included in the notebook, be sure that the Python code is successfully imported and included in your submission if necessary.

> **Note**: Once you have completed all of the code implementations, you need to finalize your work by exporting the iPython Notebook as an HTML document. Before exporting the notebook to html, all of the code cells need to have been run so that reviewers can see the final implementation and output. You can then export the notebook by using the menu above and navigating to \n", "**File -> Download as -> HTML (.html)**. Include the finished document along with this notebook as your submission.

In addition to implementing code, there is a writeup to complete. The writeup should be completed in a separate file, which can be either a markdown file or a pdf document. There is a write up template (https://github.com/udacity/CarND-Traffic-Sign-Classifier-Project/blob/master/writeup_template.md) that can be used to guide the writing process. Completing the code template and writeup template will cover all of the rubric points (https://review.udacity.com/#!/rubrics/481/view) for this project.

The rubric (https://review.udacity.com/#!/rubrics/481/view) contains "Stand Out Suggestions" for enhancing the project beyond the minimum requirements. The stand out suggestions are optional. If you decide to pursue the "stand out suggestions", you can include the code in this Ipython notebook and also discuss the results in the writeup file.

> **Note:** Code and Markdown cells can be executed using the **Shift + Enter** keyboard shortcut. In addition, Markdown cells can be edited by typically double-clicking the cell to enter edit mode.

Step -1: Preparations

Launching on AWS EC2 and other preparations

- go to console (https://us-east-2.console.aws.amazon.com/ec2/v2/home?region=us-east-2)
- Click on "Launch instance"
- Select community and select "udacity-carnd"
- Filter gy "GPU graphics"
- Edit security group. Add rule to allow connections to port 8888. Type: "Custom TCP Rule", Protocol "TCP", Port Range "8888", Source "Anywhere" (leave the "0.0.0.0/0")
- Edit security group to enable SSH (add SSH rule)
- Click launch instance, select "Proceed without a key pair" and acknowledge that you know the password (carnd) and user (carnd)

- Wait for instance to pass. Copy IP-Adress
- ssh instance
- `source activate carnd-term1`
- Clone LeNetLab `git clone https://github.com/BernhardSchlegel/CarND-Traffic-Sign-Classifier-Project`
- cd into folder, `mkdir traffic-signs-data` and `cd traffic-signs-data`
- Download and unzip traffic sign dataset `wget https://d17h27t6h515a5.cloudfront.net/topher/2017/February/5898cd6f_traffic-signs-data/traffic-signs-data.zip`, followed by `unzip traffic-signs-data.zip` or click [here (https://d17h27t6h515a5.cloudfront.net/topher/2017/February/5898cd6f_traffic-signs-data/traffic-signs-data.zip)](https://d17h27t6h515a5.cloudfront.net/topher/2017/February/5898cd6f_traffic-signs-data/traffic-signs-data.zip).
- Move folder up `cd ..` and start jupyter notebook `jupyter notebook`

If you're on a GPU system

- Install GPU support as described [here (https://www.tensorflow.org/install/#optional_install_cuda_gpus_on_linux)](https://www.tensorflow.org/install/#optional_install_cuda_gpus_on_linux)
- Execute `pip install tensorflow-gpu`

Installing OpenCV

easy `conda install --channel https://conda.anaconda.org/menpo opencv3`

hard:

```
sudo apt-get -y update
sudo apt-get -y upgrade
sudo apt-get -y install build-essential cmake git pkg-config
sudo apt-get -y install libjpeg8-dev libtiff5-dev libjasper-dev libpng12-de
v libavcodec-dev libavformat-dev libswscale-dev libv4l-dev
sudo apt-get -y install libgtk2.0-dev
sudo apt-get -y install libatlas-base-dev gfortran
cd ~
git clone https://github.com/Itseez/opencv.git
cd opencv
git checkout 3.2.0
cd ~
git clone https://github.com/Itseez/opencv_contrib.git
cd opencv_contrib
git checkout 3.2.0
cd ~/opencv
mkdir build
cd build
cmake -D CMAKE_BUILD_TYPE=RELEASE \
    -D CMAKE_INSTALL_PREFIX=/usr/local \
    -D INSTALL_C_EXAMPLES=ON \
    -D INSTALL_PYTHON_EXAMPLES=ON \
    -D OPENCV_EXTRA_MODULES_PATH=~/opencv_contrib/modules \
    -D BUILD_EXAMPLES=ON ..
make -j8
sudo make install
sudo ldconfig
cd ~/.virtualenvs/cv/lib/python3.5/site-packages/
ln -s /usr/local/lib/python3.5/site-packages/cv2.cpython-34m.so cv2.so
```

In [1]:

```python
import time
# define Log function
def log(text):
    print(time.strftime('%Y.%m.%d, %H:%M:%S') + ': ' + text)


def ping():
    return datetime.datetime.now()


def pong(dt):
    now = datetime.datetime.now()
    diff = now - dt
    ms = round(diff.total_seconds()*1000)
    return ms
```

Step 0: Load The Data

In [2]:

```python
import os

# to save time after notebook reset, save and load data
def save_data():

    directory = "processed"

    if not os.path.exists(directory):
        os.makedirs(directory)

    np.save("processed/X_train", X_train)
    np.save("processed/y_train", y_train)
    np.save("processed/X_valid", X_valid)
    np.save("processed/y_valid", y_valid)
    np.save("processed/X_test_prep", X_test_prep)
    np.save("processed/y_test", y_test)

def load_data():
    global X_train
    global y_train
    global X_valid
    global y_valid
    global X_test_prep
    global y_test

    directory = "processed"

    if os.path.exists(directory):
        X_train = np.load("processed/X_train.npy")
        y_train = np.load("processed/y_train.npy")
        X_valid = np.load("processed/X_valid.npy")
        y_valid = np.load("processed/y_valid.npy")
        X_test_prep = np.load("processed/X_test_prep.npy")
        y_test = np.load("processed/y_test.npy")
```

In [3]:

```python
# Load pickled data
import pickle

training_file =  'traffic-signs-data/train.p'
validation_file= 'traffic-signs-data/valid.p'
testing_file =   'traffic-signs-data/test.p'

with open(training_file, mode='rb') as f:
    train = pickle.load(f)
with open(validation_file, mode='rb') as f:
    valid = pickle.load(f)
with open(testing_file, mode='rb') as f:
    test = pickle.load(f)

X_train, y_train = train['features'], train['labels']
X_valid, y_valid = valid['features'], valid['labels']
X_test, y_test = test['features'], test['labels']
```

Step 1: Dataset Summary & Exploration

The pickled data is a dictionary with 4 key/value pairs:

- `'features'` is a 4D array containing raw pixel data of the traffic sign images, (num examples, width, height, channels).
- `'labels'` is a 1D array containing the label/class id of the traffic sign. The file `signnames.csv` contains id -> name mappings for each id.
- `'sizes'` is a list containing tuples, (width, height) representing the original width and height the image.
- `'coords'` is a list containing tuples, (x1, y1, x2, y2) representing coordinates of a bounding box around the sign in the image. **THESE COORDINATES ASSUME THE ORIGINAL IMAGE. THE PICKLED DATA CONTAINS RESIZED VERSIONS (32 by 32) OF THESE IMAGES**

Complete the basic data summary below. Use python, numpy and/or pandas methods to calculate the data summary rather than hard coding the results. For example, the pandas shape method (http://pandas.pydata.org/pandas-docs/stable/generated/pandas.DataFrame.shape.html) might be useful for calculating some of the summary results.

Provide a Basic Summary of the Data Set Using Python, Numpy and/or Pandas

In [4]:

```python
import numpy as np
### Replace each question mark with the appropriate value.
### Use python, pandas or numpy methods rather than hard coding the results

# TODO: Number of training examples
n_train = len(X_train)

# TODO: Number of validation examples
n_validation = len(X_valid)

# TODO: Number of testing examples.
n_test = len(X_test)

# TODO: What's the shape of an traffic sign image?
image_shape = X_train[0].shape

# TODO: How many unique classes/labels there are in the dataset.
classes, class_indices, class_counts = np.unique(y_train, return_index = True, retu
n_classes = len(classes)
log("Number of training examples ={}".format(n_train))
log("Number of testing examples ={}".format(n_test))
log("Number of validation examples ={}".format(n_validation))
log("Image data shape ={}".format(image_shape))
log("Number of classes ={}".format(n_classes))

assert(len(X_train) == len(y_train))
assert(len(X_valid) == len(y_valid))
assert(len(X_test) == len(y_test))
```

```
2017.09.01, 14:19:55: Number of training examples =34799
2017.09.01, 14:19:55: Number of testing examples =12630
2017.09.01, 14:19:55: Number of validation examples =4410
2017.09.01, 14:19:55: Image data shape =(32, 32, 3)
2017.09.01, 14:19:55: Number of classes =43
```

Include an exploratory visualization of the dataset

Visualize the German Traffic Signs Dataset using the pickled file(s). This is open ended, suggestions include: plotting traffic sign images, plotting the count of each sign, etc.

The Matplotlib (http://matplotlib.org/) examples (http://matplotlib.org/examples/index.html) and gallery (http://matplotlib.org/gallery.html) pages are a great resource for doing visualizations in Python.

NOTE: It's recommended you start with something simple first. If you wish to do more, come back to it after you've completed the rest of the sections. It can be interesting to look at the distribution of classes in the training, validation and test set. Is the distribution the same? Are there more examples of some classes than others?

In [5]:

```python
import pandas as pd
signnames = pd.DataFrame.from_csv('signnames.csv', index_col=None)
def get_sign_name(sign_number):
    return signnames[signnames['ClassId'] == sign_number].iloc[0,1]
```

In [6]:

```python
### Data exploration visualization code goes here.
### Feel free to use as many code cells as needed.
import matplotlib.pyplot as plt
# Visualizations will be shown in the notebook.
%matplotlib inline
```

In [7]:

```python
# plot class distributio
def plot_class_counts(class_counts, heading = "dataset"):
    plt.bar( np.arange( 43 ), class_counts, align='center' )
    plt.xlabel('Class')
    #plt.xlabel(np.array(signnames['SignName']), rotation = 90)
    plt.ylabel('Number of training examples')
    plt.xlim([-1, 43])
    plt.title(heading)
    plt.show()

plot_class_counts(class_counts, heading = "dataset")
```

In [8]:

```python
import random
import numpy as np
import matplotlib.pyplot as plt
%matplotlib inline

def plot_image(img):
    plt.figure(figsize=(1,1))
    plt.imshow(img.squeeze(), cmap="gray")

idx = None # ensure same images to show after each step
def plot_images_for_class(X_train, y_train, mode = 'random', max_classes = 3, n_per
    global idx
    global n_classes

    for i in classes:
        if i >= max_classes:
            break

        log('-----------------------------------------------------------------')
        log('Examples for class {}, \"{}\" ({} samples)'.format(i, get_sign_name(i)

        # get 5 random indices
        X_train_class = X_train[y_train == i]

        if mode is 'random':
            if idx is None:
                idx = np.random.randint(X_train_class.shape[0], size=n_per_class)
            X_sub = X_train_class[idx]
        elif mode is 'deterministic':
            if X_train_class.shape[0] < n_per_class:
                n_classes = X_train_class.shape[0]
            X_sub = X_train_class[0:n_per_class]

        fig, axs = plt.subplots(nrows=1, ncols=n_per_class)

        for j in range(0,n_per_class):
            image = X_sub[j].squeeze()

            axs[j].imshow(image, cmap="gray")

        plt.show()
```

In [9]:

```
plot_images_for_class(X_train, y_train, max_classes = 3)
```


2017.09.01, 14:20:08: --

2017.09.01, 14:20:08: Examples for class 1, "Speed limit (30km/h)"
(1980 samples)

2017.09.01, 14:20:08: --

2017.09.01, 14:20:08: Examples for class 2, "Speed limit (50km/h)"
(2010 samples)

In [10]:

```
# image_num, flip_x, flip_y, flip_xy, new class after transformation (-1 new = old)
augmentables = [
    [0, 0, 0, 0, -1],    # class Speed limit (20km/h)
    [1, 0, 0, 0, -1],    # class Speed limit (30km/h)
    [2, 0, 0, 0, -1],    # class Speed limit (50km/h)
    [3, 0, 0, 0, -1],    # class Speed limit (60km/h)
    [4, 0, 0, 0, -1],    # class Speed limit (70km/h)
    [5, 0, 0, 0, -1],    # class Speed limit (80km/h)
    [6, 0, 0, 0, -1],    # class End of speed limit (80km/h)
    [7, 0, 0, 0, -1],    # class Speed limit (100km/h)
    [8, 0, 0, 0, -1],    # class Speed limit (120km/h)
    [9, 1, 0, 0, -1],    # class No passing
    [10, 0, 0, 0, -1],   # class No passing for vehicles over 3.5 metric tons
    [11, 1, 0, 0, -1],   # class Right-of-way at the next intersection
    [12, 1, 1, 1, -1],   # class Priority road
    [13, 1, 0, 0, -1],   # class Yield
    [14, 0, 0, 0, -1],   # class Stop
    [15, 1, 1, 1, -1],   # class No vehicles
    [16, 0, 0, 0, -1],   # class Vehicles over 3.5 metric tons prohibited
    [17, 1, 1, 1, -1],   # class No entry
    [18, 1, 0, 0, -1],   # class General caution
    [19, 1, 0, 0, 20],   # class Dangerous curve to the left
    [20, 1, 0, 0, 19],   # class Dangerous curve to the right
    [21, 0, 0, 0, -1],   # class Double curve
    [22, 1, 0, 0, -1],   # class Bumpy road
    [23, 0, 0, 0, -1],   # class Slippery road
    [24, 0, 0, 0, -1],   # class Road narrows on the right
    [25, 0, 0, 0, -1],   # class Road work
    [26, 1, 0, 0, -1],   # class Traffic signals
    [27, 0, 0, 0, -1],   # class Pedestrians
    [28, 0, 0, 0, -1],   # class Children crossing
    [29, 0, 0, 0, -1],   # class Bicycles crossing
    [30, 1, 0, 0, -1],   # class Beware of ice/snow
    [31, 0, 0, 0, -1],   # class Wild animals crossing
    [32, 0, 0, 1, -1],   # class End of all speed and passing limits
    [33, 1, 0, 0, 34],   # class Turn right ahead
    [34, 1, 0, 0, 33],   # class Turn left ahead
    [35, 1, 0, 0, -1],   # class Ahead only
    [36, 1, 0, 0, 37],   # class Go straight or right
    [37, 1, 0, 0, 36],   # class Go straight or left
    [38, 1, 0, 0, 39],   # class Keep right
    [39, 1, 0, 0, 38],   # class Keep left
    [40, 0, 0, 0, -1],   # class Roundabout mandatory
    [41, 0, 0, 0, -1],   # class End of no passing
    [42, 0, 0, 0, -1]    # class End of no passing by vehicles over 3.5 metric tons
]
```

In [11]:

```python
from skimage.transform import warp
from skimage.transform import rotate
from skimage.transform import rescale
from skimage.transform import ProjectiveTransform
import cv2

def flip_X(img):
    return img.copy()[:, ::-1]

def flip_Y(img):
    return img.copy()[::-1, :]

def augment_by_flipping(X, y):
    # augment as much as possible
    X_augmented = []
    y_augmented = []

    for image_num, flip_x, flip_y, flip_xy, new_class in augmentables:
        X_sub_image = X[y == image_num]
        n_images_of_class = X_sub_image.shape[0]
        if (flip_x + flip_y + flip_xy) > 0:
            log("found {} images of class {}. Applying: flip_x={}, flip_y={}, flip_
                n_images_of_class, image_num, flip_x, flip_y, flip_xy))

            counter = 0
            for image in X_sub_image:
                if new_class is -1:
                    new_class = image_num

                if flip_x > 0:
                    X_augmented.append(flip_X(image))
                    y_augmented.append(new_class)
                    counter += 1

                if flip_y > 0:
                    X_augmented.append(flip_Y(image))
                    y_augmented.append(new_class)
                    counter += 1

                if flip_xy > 0:
                    image = flip_Y(image)
                    X_augmented.append(flip_X(image))
                    y_augmented.append(new_class)
                    counter += 1

            log("created {} artificial samples.".format(counter))
        else:
            log("class {} not transformable".format(image_num))

    return np.array(X_augmented), y_augmented


def modify_warp(img, intensity = 0.25, depth_layer = 0):
    rows,cols,ch = img.shape

    assert len(img.shape) is 3, "this method only works for images with a depth cha
    assert ch is 1, "this method only works for images with one channel"
```

```python
    img_shape_x, img_shape_y, depth = img.shape
    p1_left_x = int(img_shape_x * intensity)
    p1_right_x = int(img_shape_x * ( 1 - intensity))
    p1_top_y = int(img_shape_y * intensity)
    p1_bot_y = int(img_shape_y * ( 1 - intensity))
    p2_left_x = int(img_shape_x * intensity)
    p2_right_x = int(img_shape_x * ( 1 - intensity))
    p2_top_y = int(img_shape_y * intensity)
    p2_bot_y = int(img_shape_y * ( 1 - intensity))

    # points in array are numbered like follows
    #  1. 3.
    #  2. 4.
    pts1 = np.float32([[p1_left_x, p1_top_y], [p1_left_x, p1_bot_y], [p1_right_x, p
    pts2 = np.float32([[p1_left_x - random.uniform(0, p1_left_x) , p1_top_y - rando
                       [p1_left_x - random.uniform(0, p1_left_x) , p1_bot_y + rando
                       [p1_right_x + random.uniform(0, img_shape_x - p1_right_x) ,
                       [p1_right_x + random.uniform(0, img_shape_x - p1_right_x), p

    M = cv2.getPerspectiveTransform(pts1,pts2)
    dst = np.ones(shape = [32,32,1])
    dst[:,:,depth_layer] = np.array(cv2.warpPerspective(img,M,(img_shape_y,img_shap

    return dst


def augment_by_transforming(X, y, class_counts, intensity = 0.2, min_multiplier = 1
    """
    This function will modify each class count so that every class is present
    number of times of the most present class times the min_multiplier
    """

    X_augmented = []
    y_augmented = []

    n_idx_max = np.argmax(class_counts)
    n_max = class_counts[n_idx_max]
    target_samples = n_max * min_multiplier
    log("class {} ({}) is the maximally represented class (samples={}). Target samp
        n_idx_max, get_sign_name(n_idx_max), n_max, target_samples))

    for class_num in range(len(class_counts)):
        X_sub_image = X[y == class_num]
        n_class = class_counts[class_num]
        oversampling_ratio = target_samples / n_class
        samples_to_create = int(oversampling_ratio * n_class) - n_class
        log("class {0} is present {1} times. Oversampling ratio is {2:.3}. Samples
                class_num, n_class, oversampling_ratio, samples_to_create))

        n_generated = 0
        while n_generated < samples_to_create:
            # iterate through class instances
            image = X_sub_image[n_generated%class_counts[class_num]]

            # add to arrray (X and y)
            X_augmented.append(modify_warp(image, intensity = intensity))
            y_augmented.append(class_num)

            n_generated += 1
            if int(n_generated % (samples_to_create / 10)) is 0 and debug is True:
                log('generating {} / {} sample. '.format(n_generated, samples_to_cr
```

```
        #else:
        #   log('## {} / {}'.format(n_generated, samples_to_create))

    return np.array(X_augmented), y_augmented
```

In [12]:

```python
def augment_dataset(X, y):
    # before
    classes, class_indices, class_counts = np.unique(y, return_index = True, return
    plot_class_counts(class_counts, "dataset original")

    # augment by flipping
    X_flipped, y_flipped = augment_by_flipping(X, y)
    log("created {} new samples by flipping.".format(X_flipped.shape[0]))
    X_train_augmented = np.append(X, X_flipped, axis = 0)
    y_train_augmented = np.append(y, y_flipped, axis = 0)
    classes, class_indices, class_counts = np.unique(y_train_augmented, return_inde
    plot_class_counts(class_counts)

    # augment by transforming
    X_transformed, y_transformed = augment_by_transforming(X_train_augmented, y_tra
    log("created {} new samples by transforming.".format(X_transformed.shape[0] - X
    X_train_augmented = np.append(X_train_augmented, X_transformed, axis = 0)
    y_train_augmented = np.append(y_train_augmented, y_transformed, axis = 0)
    classes, class_indices, class_counts = np.unique(y_train_augmented, return_inde
    plot_class_counts(class_counts, "dataset original")

    return X_train_augmented, y_train_augmented
```

Step 2: Design and Test a Model Architecture

Design and implement a deep learning model that learns to recognize traffic signs. Train and test your model on the German Traffic Sign Dataset (http://benchmark.ini.rub.de/?section=gtsrb&subsection=dataset).

The LeNet-5 implementation shown in the classroom (https://classroom.udacity.com/nanodegrees/nd013/parts/fbf77062-5703-404e-b60c-95b78b2f3f9e/modules/6df7ae49-c61c-4bb2-a23e-6527e69209ec/lessons/601ae704-1035-4287-8b11-e2c2716217ad/concepts/d4aca031-508f-4e0b-b493-e7b706120f81) at the end of the CNN lesson is a solid starting point. You'll have to change the number of classes and possibly the preprocessing, but aside from that it's plug and play!

With the LeNet-5 solution from the lecture, you should expect a validation set accuracy of about 0.89. To meet specifications, the validation set accuracy will need to be at least 0.93. It is possible to get an even higher accuracy, but 0.93 is the minimum for a successful project submission.

There are various aspects to consider when thinking about this problem:

- Neural network architecture (is the network over or underfitting?)
- Play around preprocessing techniques (normalization, rgb to grayscale, etc)
- Number of examples per label (some have more than others).
- Generate fake data.

Here is an example of a published baseline model on this problem (http://yann.lecun.com/exdb/publis/pdf/sermanet-ijcnn-11.pdf). It's not required to be familiar with the approach used in the paper but, it's good practice to try to read papers like these.

Pre-process the Data Set (normalization, grayscale, etc.)

Minimally, the image data should be normalized so that the data has mean zero and equal variance. For image data, (`pixel - 128)/ 128` is a quick way to approximately normalize the data and can be used in this project.

Other pre-processing steps are optional. You can try different techniques to see if it improves performance.

Use the code cell (or multiple code cells, if necessary) to implement the first step of your project.

In [13]:

```python
from skimage import exposure
import sklearn.preprocessing as pp

def rgb2gray(rgb):
    new = np.empty([rgb.shape[0], rgb.shape[1], rgb.shape[2], 1])
    new[:,:,:,0] = np.dot(rgb[...,:3], [0.299, 0.587, 0.114])
    return new
```

In [14]:

```python
def preprocess_images(X, do_hist = True):
    # grayscale
    X = rgb2gray(X)

    # normalize
    X =  X = (X / 255.).astype(np.float32)

    # histogram localication
    if do_hist:
        for i in range(X.shape[0]):
            X[i,:,:,0] = exposure.equalize_adapthist(X[i,:,:,0])

    return X
```

In [15]:

```python
log('processing training')
X_train = preprocess_images(X_train, do_hist=True)
```

```
2017.09.01, 14:20:19: processing training

/home/q372283/anaconda3/lib/python3.5/site-packages/skimage/util/dtyp
e.py:122: UserWarning: Possible precision loss when converting from fl
oat32 to uint16
  .format(dtypeobj_in, dtypeobj_out))
```

In [16]:

```
log('processing validation')
X_valid = preprocess_images(X_valid, do_hist=True)
```

2017.09.01, 14:24:57: processing validation

/home/q372283/anaconda3/lib/python3.5/site-packages/skimage/util/dtyp
e.py:122: UserWarning: Possible precision loss when converting from fl
oat32 to uint16
  .format(dtypeobj_in, dtypeobj_out))

In [17]:

```
log('augmenting train')
X_train_aug, y_train_aug = augment_dataset(X_train, y_train)
```

2017.09.01, 14:25:32: found 1170 images of class 11. Applying: flip_
x=1, flip_y=0, flip_xy=0
2017.09.01, 14:25:32: created 1170 artificial samples.
2017.09.01, 14:25:32: found 1890 images of class 12. Applying: flip_
x=1, flip_y=1, flip_xy=1
2017.09.01, 14:25:32: created 5670 artificial samples.
2017.09.01, 14:25:32: found 1920 images of class 13. Applying: flip_
x=1, flip_y=0, flip_xy=0
2017.09.01, 14:25:32: created 1920 artificial samples.
2017.09.01, 14:25:32: class 14 not transformable
2017.09.01, 14:25:32: found 540 images of class 15. Applying: flip_x
=1, flip_y=1, flip_xy=1
2017.09.01, 14:25:32: created 1620 artificial samples.
2017.09.01, 14:25:32: class 16 not transformable
2017.09.01, 14:25:32: found 990 images of class 17. Applying: flip_x
=1, flip_y=1, flip_xy=1
2017.09.01, 14:25:32: created 2970 artificial samples.
2017.09.01, 14:25:32: found 1080 images of class 18. Applying: flip_
x=1, flip_y=0, flip_xy=0
2017.09.01, 14:25:32: created 1080 artificial samples.

In [ ]:

```
X_test_prep = preprocess_images(X_test)
```

In [18]:

```
log("training with a total number of {} samples".format(y_train.shape[0]))
```

2017.09.01, 14:25:54: training with a total number of 34799 samples

In [19]:

```
idx = None
plot_images_for_class(X_train_aug, y_train_aug, max_classes=3)
```

2017.09.01, 14:25:54: --

2017.09.01, 14:25:54: Examples for class 0, "Speed limit (20km/h)" (18
0 samples)

2017.09.01, 14:25:55: --

2017.09.01, 14:25:55: Examples for class 1, "Speed limit (30km/h)" (19
80 samples)

2017.09.01, 14:25:55: --

2017.09.01, 14:25:55: Examples for class 2, "Speed limit (50km/h)" (20
10 samples)

In [20]:

```
from sklearn.utils import shuffle
```

In [21]:

```
X_train, y_train = shuffle(X_train_aug, y_train_aug)
```

In [22]:

```
log('augmenting valid')
X_valid_aug, y_valid_aug = augment_dataset(X_valid, y_valid)
```

In [23]:

```
X_valid, y_valid = X_valid_aug, y_valid_aug
```

In [ ]:

```
save_data()
```

Model Architecture

In [24]:

```
import tensorflow as tf

n_classes = 43
image_depth = 1
```

In [25]:

```python
from tensorflow.contrib.layers import flatten

def conv2d(input, filter_width, filter_height, stage_name = "default", input_depth
                  stride = 2):
    # This solution uses the tf.truncated_normal() function to initialize the weigh
    # Using the default mean and standard deviation from tf.truncated_normal() is f
    # these hyperparameters can result in better performance.
    mu = 0
    sigma = 0.1

    # Filter (weights and bias)
    # The shape of the filter weight is (height, width, input_depth, output_depth)
    # The shape of the filter bias is (output_depth,)
    # TODO: Define the filter weights `F_W` and filter bias `F_b`.
    # NOTE: Remember to wrap them in ` tf.Variable`, they are trainable parameters
    F_W = tf.Variable(tf.truncated_normal(shape = (filter_width, filter_height, inp
                                          mean = mu,
                                          stddev = sigma), name = "weight_" + stage_
    tf.add_to_collection(tf.GraphKeys.REGULARIZATION_LOSSES, F_W)
    F_b = tf.Variable(tf.zeros(output_depth), name = "bias_" + stage_name)
    # TODO: Set the stride for each dimension (batch_size, height, width, depth)
    strides = [1, stride, stride, 1]
    # TODO: set the padding, either 'VALID' or 'SAME'.
    padding = 'VALID'
    # https://www.tensorflow.org/versions/r0.11/api_docs/python/nn.html#conv2d
    # `tf.nn.conv2d` does not include the bias computation so we have to add it our
    return tf.nn.conv2d(input, F_W, strides, padding) + F_b

def maxpool(input, filter_height, filter_width, stride):
    # TODO: Set the ksize (filter size) for each dimension (batch_size, height, wid
    ksize = [1, filter_height, filter_width, 1]
    # TODO: Set the stride for each dimension (batch_size, height, width, depth)
    strides = [1, stride, stride, 1]
    # TODO: set the padding, either 'VALID' or 'SAME'.
    padding = 'VALID'
    # https://www.tensorflow.org/versions/r0.11/api_docs/python/nn.html#max_pool
    return tf.nn.max_pool(input, ksize, strides, padding)


def LeNet3(x, keep_prob = tf.placeholder(tf.float32), n_classes = 43, original_dept

    # Arguments used for tf.truncated_normal, randomly defines variables for the we
    mu = 0
    sigma = 0.1

    # TODO: Layer 1: Convolutional. Input = 32x32x1. Output = 28x28x6.
    layer1_conv = conv2d(x, 5, 5, stage_name = "layer1_conv", input_depth = origina

    # TODO: Activation.
    # Use predefined relu function as activation
    layer1_relu = tf.nn.relu(layer1_conv)

    # TODO: Pooling. Input = 28x28x6. Output = 14x14x6.
    layer1_pooled = maxpool(layer1_relu, filter_height = 2, filter_width = 2, strid

    # TODO: Layer 2: Convolutional. Output = 10x10x16.
    layer2_conv = conv2d(layer1_pooled, filter_height = 5, stage_name = "layer2_con

    # TODO: Activation.
```

```
    layer2_relu = tf.nn.relu(layer2_conv)

    # UPGRADE: Dropout
    layer2_relu = tf.nn.dropout(layer2_relu, keep_prob)

    # TODO: Pooling. Input = 10x10x16. Output = 5x5x16.
    layer2_pooled = maxpool(layer2_relu, filter_height = 2, filter_width = 2, strid

    # TODO: Flatten. Input = 5x5x16. Output = 400.
    flattened = tf.contrib.layers.flatten(layer2_pooled)

    # TODO: Layer 3: Fully Connected. Input = 400. Output = 120.
    layer3 = tf.contrib.layers.fully_connected(flattened, 120)

    # TODO: Activation.
    layer3_relu = tf.nn.relu(layer3)

    # UPGRADE: Dropout
    layer3_relu = tf.nn.dropout(layer3_relu, keep_prob)

    # SOLUTION: Layer 4: Fully Connected. Input = 120. Output = 84.
    fc2_W  = tf.Variable(tf.truncated_normal(shape=(120, 84), mean = mu, stddev = s
    fc2_b  = tf.Variable(tf.zeros(84), name= 'bias_fc2_b')
    layer4     = tf.matmul(layer3_relu, fc2_W) + fc2_b

    # TODO: Activation.
    layer4_relu = tf.nn.relu(layer4)

    # SOLUTION: Layer 5: Fully Connected. Input = 84. Output = 10.
    fc3_W  = tf.Variable(tf.truncated_normal(shape=(84, n_classes), mean = mu, stdd
    fc3_b  = tf.Variable(tf.zeros(n_classes), name= 'bias_fc3_b')
    logits = tf.matmul(layer4_relu, fc3_W) + fc3_b

    return logits
```

In [26]:

```
x = tf.placeholder(tf.float32, (None, 32, 32, image_depth))
y = tf.placeholder(tf.int32, (None))
one_hot_y = tf.one_hot(y, n_classes)
```

Train, Validate and Test the Model

A validation set can be used to assess how well the model is performing. A low accuracy on the training and validation sets imply underfitting. A high accuracy on the training set but low accuracy on the validation set implies overfitting.

Train

In [7]:

```
load_data()
```

In [27]:

```python
rate_coarse = 0.0005
rate_fine = 0.0001
EPOCHS = 30
BATCH_SIZE = 256
BETA = 0.001 # This is a good beta value to start with for regularization

keep_prob = tf.placeholder(tf.float32)
logits = LeNet3(x, keep_prob, n_classes= n_classes, original_depth= image_depth)
cross_entropy = tf.nn.softmax_cross_entropy_with_logits(labels=one_hot_y, logits=lo
loss_operation = tf.reduce_mean(cross_entropy)

# EXPERIMENTAL
# Loss function using L2 Regularization
trainable_vars   = tf.trainable_variables()
lossL2 = tf.add_n([ tf.nn.l2_loss(v) for v in trainable_vars
                    if 'bias' not in v.name ])
#weights = tf.get_collection(tf.GraphKeys.REGULARIZATION_LOSSES)
#regularizer = tf.nn.l2_loss(weights)
loss_operation = tf.reduce_mean(loss_operation + BETA * lossL2)
# EXPERIMENTAL

optimizer = tf.train.AdamOptimizer(learning_rate = rate_coarse)
training_operation = optimizer.minimize(loss_operation)
```

In [28]:

```python
correct_prediction = tf.equal(tf.argmax(logits, 1), tf.argmax(one_hot_y, 1))
accuracy_operation = tf.reduce_mean(tf.cast(correct_prediction, tf.float32))
saver = tf.train.Saver()

def evaluate(X_data, y_data):
    num_examples = len(X_data)
    total_accuracy = 0
    sess = tf.get_default_session()
    for offset in range(0, num_examples, BATCH_SIZE):
        batch_x, batch_y = X_data[offset:offset+BATCH_SIZE], y_data[offset:offset+B
        accuracy = sess.run(accuracy_operation, feed_dict={x: batch_x, y: batch_y,
        total_accuracy += (accuracy * len(batch_x))
    return total_accuracy / num_examples
```

In [29]:

```python
def train_net(sess, X_train, y_train):
    accuracies = np.zeros(n_auto_stop).tolist()

    sess.run(tf.global_variables_initializer())
    num_examples = len(X_train)

    log("Training...")
    for i in range(EPOCHS):
        X_train, y_train = shuffle(X_train, y_train)
        for offset in range(0, num_examples, BATCH_SIZE):
            end = offset + BATCH_SIZE
            batch_x, batch_y = X_train[offset:end], y_train[offset:end]
            sess.run(training_operation, feed_dict={x: batch_x, y: batch_y, keep_pr

        validation_accuracy = evaluate(X_valid, y_valid)
        log("EPOCH {} ...".format(i+1))
        log("Validation Accuracy = {:.3f}".format(validation_accuracy))
        if len(accuracies) > n_auto_stop:
            accuracies = accuracies[-n_auto_stop:] # take last n_auto_stop elements
        if min(accuracies) >= (validation_accuracy + 0.1):
            log("accuracy stopped increasing more than 0.1. stopping...")
            break
        else:
            log("minimum accuracy during last {0} epochs was {1:.3}, tends to incre
                n_auto_stop, min(accuracies)))

        accuracies.append(validation_accuracy)

    saver.save(sess, './lenet')
    log("Model saved")
```

In [33]:

```
EPOCHS = 10
n_auto_stop = 10 # determines how many epochs are used for auto stopping, this is i
with tf.Session() as sess:
    train_net(sess, X_train, y_train)
```

```
2017.09.01, 14:42:17: Training...
2017.09.01, 14:42:36: EPOCH 1 ...
2017.09.01, 14:42:36: Validation Accuracy = 0.874
2017.09.01, 14:42:36: minimum accuracy during last 10 epochs was 0.0,
 tends to increase. Continueing...
2017.09.01, 14:42:56: EPOCH 2 ...
2017.09.01, 14:42:56: Validation Accuracy = 0.917
2017.09.01, 14:42:56: minimum accuracy during last 10 epochs was 0.0,
 tends to increase. Continueing...
2017.09.01, 14:43:16: EPOCH 3 ...
2017.09.01, 14:43:16: Validation Accuracy = 0.932
2017.09.01, 14:43:16: minimum accuracy during last 10 epochs was 0.0,
 tends to increase. Continueing...
2017.09.01, 14:43:37: EPOCH 4 ...
2017.09.01, 14:43:37: Validation Accuracy = 0.942
2017.09.01, 14:43:37: minimum accuracy during last 10 epochs was 0.0,
 tends to increase. Continueing...
2017.09.01, 14:43:58: EPOCH 5 ...
2017.09.01, 14:43:58: Validation Accuracy = 0.945
2017.09.01, 14:43:58: minimum accuracy during last 10 epochs was 0.0,
 tends to increase. Continueing...
2017.09.01, 14:44:19: EPOCH 6 ...
2017.09.01, 14:44:19: Validation Accuracy = 0.946
2017.09.01, 14:44:19: minimum accuracy during last 10 epochs was 0.0,
 tends to increase. Continueing...
2017.09.01, 14:44:40: EPOCH 7 ...
2017.09.01, 14:44:40: Validation Accuracy = 0.951
2017.09.01, 14:44:40: minimum accuracy during last 10 epochs was 0.0,
 tends to increase. Continueing...
2017.09.01, 14:45:00: EPOCH 8 ...
2017.09.01, 14:45:00: Validation Accuracy = 0.949
2017.09.01, 14:45:00: minimum accuracy during last 10 epochs was 0.0,
 tends to increase. Continueing...
2017.09.01, 14:45:21: EPOCH 9 ...
2017.09.01, 14:45:21: Validation Accuracy = 0.957
2017.09.01, 14:45:21: minimum accuracy during last 10 epochs was 0.0,
 tends to increase. Continueing...
2017.09.01, 14:45:42: EPOCH 10 ...
2017.09.01, 14:45:42: Validation Accuracy = 0.953
2017.09.01, 14:45:42: minimum accuracy during last 10 epochs was 0.0,
 tends to increase. Continueing...
2017.09.01, 14:45:42: Model saved
```

In [35]:

```python
# train again with finer settings
optimizer = tf.train.AdamOptimizer(learning_rate = rate_fine)
training_operation = optimizer.minimize(loss_operation)
n_auto_stop = 10
EPOCHS = 100

with tf.Session() as sess:
    saver.restore(sess, tf.train.latest_checkpoint('.'))
    train_net(sess, X_train, y_train)
```

```
INFO:tensorflow:Restoring parameters from ./lenet
2017.09.01, 14:57:45: Training...
2017.09.01, 14:58:05: EPOCH 1 ...
2017.09.01, 14:58:05: Validation Accuracy = 0.650
2017.09.01, 14:58:05: minimum accuracy during last 10 epochs was 0.0,
 tends to increase. Continueing...
2017.09.01, 14:58:26: EPOCH 2 ...
2017.09.01, 14:58:26: Validation Accuracy = 0.765
2017.09.01, 14:58:26: minimum accuracy during last 10 epochs was 0.0,
 tends to increase. Continueing...
2017.09.01, 14:58:46: EPOCH 3 ...
2017.09.01, 14:58:46: Validation Accuracy = 0.823
2017.09.01, 14:58:46: minimum accuracy during last 10 epochs was 0.0,
 tends to increase. Continueing...
2017.09.01, 14:59:07: EPOCH 4 ...
2017.09.01, 14:59:07: Validation Accuracy = 0.853
2017.09.01, 14:59:07: minimum accuracy during last 10 epochs was 0.0,
 tends to increase. Continueing...
2017.09.01, 14:59:27: EPOCH 5 ...
2017.09.01, 14:59:27: Validation Accuracy = 0.871
2017.09.01, 14:59:27: minimum accuracy during last 10 epochs was 0.0,
 tends to increase. Continueing...
2017.09.01, 14:59:48: EPOCH 6 ...
2017.09.01, 14:59:48: Validation Accuracy = 0.880
2017.09.01, 14:59:48: minimum accuracy during last 10 epochs was 0.0,
 tends to increase. Continueing...
2017.09.01, 15:00:09: EPOCH 7 ...
2017.09.01, 15:00:09: Validation Accuracy = 0.893
2017.09.01, 15:00:09: minimum accuracy during last 10 epochs was 0.0,
 tends to increase. Continueing...
2017.09.01, 15:00:29: EPOCH 8 ...
2017.09.01, 15:00:29: Validation Accuracy = 0.898
2017.09.01, 15:00:29: minimum accuracy during last 10 epochs was 0.0,
 tends to increase. Continueing...
2017.09.01, 15:00:50: EPOCH 9 ...
2017.09.01, 15:00:50: Validation Accuracy = 0.908
2017.09.01, 15:00:50: minimum accuracy during last 10 epochs was 0.0,
 tends to increase. Continueing...
2017.09.01, 15:01:11: EPOCH 10 ...
2017.09.01, 15:01:11: Validation Accuracy = 0.910
2017.09.01, 15:01:11: minimum accuracy during last 10 epochs was 0.0,
 tends to increase. Continueing...
2017.09.01, 15:01:31: EPOCH 11 ...
2017.09.01, 15:01:31: Validation Accuracy = 0.919
2017.09.01, 15:01:31: minimum accuracy during last 10 epochs was 0.65,
 tends to increase. Continueing...
2017.09.01, 15:01:52: EPOCH 12 ...
2017.09.01, 15:01:52: Validation Accuracy = 0.919
2017.09.01, 15:01:52: minimum accuracy during last 10 epochs was 0.76
5, tends to increase. Continueing...
2017.09.01, 15:02:12: EPOCH 13 ...
2017.09.01, 15:02:12: Validation Accuracy = 0.928
2017.09.01, 15:02:12: minimum accuracy during last 10 epochs was 0.82
3, tends to increase. Continueing...
2017.09.01, 15:02:33: EPOCH 14 ...
2017.09.01, 15:02:33: Validation Accuracy = 0.923
2017.09.01, 15:02:33: minimum accuracy during last 10 epochs was 0.85
3, tends to increase. Continueing...
2017.09.01, 15:02:53: EPOCH 15 ...
2017.09.01, 15:02:53: Validation Accuracy = 0.929
```

2017.09.01, 15:02:53: minimum accuracy during last 10 epochs was 0.87

1, tends to increase. Continueing...
2017.09.01, 15:03:14: EPOCH 16 ...
2017.09.01, 15:03:14: Validation Accuracy = 0.935
2017.09.01, 15:03:14: minimum accuracy during last 10 epochs was 0.88,
 tends to increase. Continueing...
2017.09.01, 15:03:35: EPOCH 17 ...
2017.09.01, 15:03:35: Validation Accuracy = 0.934
2017.09.01, 15:03:35: minimum accuracy during last 10 epochs was 0.89
3, tends to increase. Continueing...
2017.09.01, 15:03:55: EPOCH 18 ...
2017.09.01, 15:03:55: Validation Accuracy = 0.936
2017.09.01, 15:03:55: minimum accuracy during last 10 epochs was 0.89
8, tends to increase. Continueing...
2017.09.01, 15:04:16: EPOCH 19 ...
2017.09.01, 15:04:16: Validation Accuracy = 0.939
2017.09.01, 15:04:16: minimum accuracy during last 10 epochs was 0.90
8, tends to increase. Continueing...
2017.09.01, 15:04:36: EPOCH 20 ...
2017.09.01, 15:04:36: Validation Accuracy = 0.937
2017.09.01, 15:04:36: minimum accuracy during last 10 epochs was 0.91,
 tends to increase. Continueing...
2017.09.01, 15:04:57: EPOCH 21 ...
2017.09.01, 15:04:57: Validation Accuracy = 0.939
2017.09.01, 15:04:57: minimum accuracy during last 10 epochs was 0.91
9, tends to increase. Continueing...
2017.09.01, 15:05:17: EPOCH 22 ...
2017.09.01, 15:05:17: Validation Accuracy = 0.941
2017.09.01, 15:05:17: minimum accuracy during last 10 epochs was 0.91
9, tends to increase. Continueing...
2017.09.01, 15:05:38: EPOCH 23 ...
2017.09.01, 15:05:38: Validation Accuracy = 0.942
2017.09.01, 15:05:38: minimum accuracy during last 10 epochs was 0.92
3, tends to increase. Continueing...
2017.09.01, 15:05:58: EPOCH 24 ...
2017.09.01, 15:05:58: Validation Accuracy = 0.943
2017.09.01, 15:05:58: minimum accuracy during last 10 epochs was 0.92
3, tends to increase. Continueing...
2017.09.01, 15:06:19: EPOCH 25 ...
2017.09.01, 15:06:19: Validation Accuracy = 0.944
2017.09.01, 15:06:19: minimum accuracy during last 10 epochs was 0.92
9, tends to increase. Continueing...
2017.09.01, 15:06:39: EPOCH 26 ...
2017.09.01, 15:06:39: Validation Accuracy = 0.945
2017.09.01, 15:06:39: minimum accuracy during last 10 epochs was 0.93
4, tends to increase. Continueing...
2017.09.01, 15:07:00: EPOCH 27 ...
2017.09.01, 15:07:00: Validation Accuracy = 0.947
2017.09.01, 15:07:00: minimum accuracy during last 10 epochs was 0.93
4, tends to increase. Continueing...
2017.09.01, 15:07:21: EPOCH 28 ...
2017.09.01, 15:07:21: Validation Accuracy = 0.948
2017.09.01, 15:07:21: minimum accuracy during last 10 epochs was 0.93
6, tends to increase. Continueing...
2017.09.01, 15:07:42: EPOCH 29 ...
2017.09.01, 15:07:42: Validation Accuracy = 0.945
2017.09.01, 15:07:42: minimum accuracy during last 10 epochs was 0.93
7, tends to increase. Continueing...
2017.09.01, 15:08:02: EPOCH 30 ...
2017.09.01, 15:08:02: Validation Accuracy = 0.948
2017.09.01, 15:08:02: minimum accuracy during last 10 epochs was 0.93

7, tends to increase. Continueing...
2017.09.01, 15:08:02: Model saved

Test

In [36]:

```
#tf.reset_default_graph()
sess = tf.get_default_session()

with tf.Session() as sess:
    saver.restore(sess, tf.train.latest_checkpoint('.'))

    test_accuracy = evaluate(X_valid, y_valid)
    log("Test Accuracy = {:.3f}".format(test_accuracy))
```

INFO:tensorflow:Restoring parameters from ./lenet
2017.09.01, 15:08:09: Test Accuracy = 0.948

In [37]:

```
idx = None # ensure same images to show after each step
plot_images_for_class(X_test_prep, y_test, max_classes = 3)
```

2017.09.01, 15:08:28: ---

2017.09.01, 15:08:28: Examples for class 0, "Speed limit (20km/h)" (18
0 samples)

2017.09.01, 15:08:28: ---

2017.09.01, 15:08:28: Examples for class 1, "Speed limit (30km/h)" (19
80 samples)

2017.09.01, 15:08:29: ---

2017.09.01, 15:08:29: Examples for class 2, "Speed limit (50km/h)" (20
10 samples)

In [38]:

```
with tf.Session() as sess:
    saver.restore(sess, tf.train.latest_checkpoint('.'))

    test_accuracy = evaluate(X_test_prep, y_test)
    log("Test Accuracy = {:.3f}".format(test_accuracy))
```

```
INFO:tensorflow:Restoring parameters from ./lenet
2017.09.01, 15:08:32: Test Accuracy = 0.939
```

Step 3: Test a Model on New Images

To give yourself more insight into how your model is working, download at least five pictures of German traffic signs from the web and use your model to predict the traffic sign type.

You may find `signnames.csv` useful as it contains mappings from the class id (integer) to the actual sign name.

Load and Output the Images

In [39]:

```python
import scipy

X_test_custom = []
y_test_custom = []

for image_name in [12, 15, 17, 35, 39]:
    im = scipy.ndimage.imread("test_images/" + str(image_name) + ".jpg")
    im_label = image_name

    X_test_custom.append(im)
    y_test_custom.append(im_label)


    plt.figure(figsize=(1,1))
    plt.title("Actual class " + str(im_label) + "(" + get_sign_name(im_label) + ")"
    plt.imshow(im.squeeze(), cmap="gray")

X_test_custom = np.array(X_test_custom)
y_test_custom = np.array(y_test_custom)
```

Actual class 12(Priority road)

Actual class 15(No vehicles)

Actual class 17(No entry)

Actual class 35(Ahead only)

Actual class 39(Keep left)

In [40]:

```
X_test_custom_prep = preprocess_images(X_test_custom)
```

/home/q372283/anaconda3/lib/python3.5/site-packages/skimage/util/dtyp
e.py:122: UserWarning: Possible precision loss when converting from fl
oat32 to uint16
  .format(dtypeobj_in, dtypeobj_out))

Predict the Sign Type for Each Image

In [30]:

```
### Run the predictions here and use the model to output the prediction for each im
### Make sure to pre-process the images with the same pre-processing pipeline used
### Feel free to use as many code cells as needed.
```

In [46]:

```
def evaluate_non_badge(X_data, y_data):

    with tf.Session() as sess:
        # Restore variables from disk.
        saver.restore(sess, "./lenet")
        print("Model restored.")

        classification = sess.run(tf.argmax(logits, 1), feed_dict={x: X_data, y: y_
        print(classification)
        return classification
```

In [47]:

```
predictions = evaluate_non_badge(X_test_custom_prep, y_test_custom)
```

INFO:tensorflow:Restoring parameters from ./lenet
Model restored.
[12 15 17 35 14]

In [102]:

```
def plot_random_image_from_trainset(X, y, label_to_plot, predicted_image):
    img1 = (X[y == label_to_plot][1])
    img2 = predicted_image

    # one time
    fig = plt.figure()

    # subplots
    ax1 = fig.add_subplot(2,2,1)
    ax1.imshow(img1.squeeze(), cmap="gray")
    ax2 = fig.add_subplot(2,2,2)
    ax2.imshow(img2.squeeze(), cmap="gray")

    # show
    plt.show()
```

In [103]:

```
for actual, prediction in zip(y_test_custom, predictions):
    print("Class {} was predicted as {}".format(actual, prediction))
    test_image = X_test_custom[y_test_custom==actual]
    plot_random_image_from_trainset(X_train_aug, y_train_aug, actual, test_image)
```

Class 12 was predicted as 12

Class 15 was predicted as 15

Class 17 was predicted as 17

Class 35 was predicted as 35

Class 39 was predicted as 14

Analyze Performance

In [ ]:

```
### Calculate the accuracy for these 5 new images.
### For example, if the model predicted 1 out of 5 signs correctly, it's 20% accura
```

In [78]:

```
TP = sum([ x==y for (x,y) in zip(y_test_custom, predictions)])
N = len(y_test_custom)
accuracy = TP / N
```

In [81]:

```
print("accuracy of custom signs is {}%".format(round(accuracy*1000)/10))
```

accuracy of custom signs is 80.0%

Output Top 5 Softmax Probabilities For Each Image Found on the Web

For each of the new images, print out the model's softmax probabilities to show the **certainty** of the model's predictions (limit the output to the top 5 probabilities for each image). `tf.nn.top_k` (https://www.tensorflow.org/versions/r0.12/api_docs/python/nn.html#top_k) could prove helpful here.

The example below demonstrates how tf.nn.top_k can be used to find the top k predictions for each image.

`tf.nn.top_k` will return the values and indices (class ids) of the top k predictions. So if k=3, for each sign, it'll return the 3 largest probabilities (out of a possible 43) and the correspoding class ids.

Take this numpy array as an example. The values in the array represent predictions. The array contains softmax probabilities for five candidate images with six possible classes. `tf.nn.top_k` is used to choose the three classes with the highest probability:

```
# (5, 6) array
a = np.array([[ 0.24879643,  0.07032244,  0.12641572,  0.34763842,  0.07893
497,
        0.12789202],
      [ 0.28086119,  0.27569815,  0.08594638,  0.0178669 ,  0.18063401,
        0.15899337],
      [ 0.26076848,  0.23664738,  0.08020603,  0.07001922,  0.1134371 ,
        0.23892179],
      [ 0.11943333,  0.29198961,  0.02605103,  0.26234032,  0.1351348 ,
        0.16505091],
      [ 0.09561176,  0.34396535,  0.0643941 ,  0.16240774,  0.24206137,
        0.09155967]])
```

Running it through `sess.run(tf.nn.top_k(tf.constant(a), k=3))` produces:

```
TopKV2(values=array([[ 0.34763842,  0.24879643,  0.12789202],
       [ 0.28086119,  0.27569815,  0.18063401],
       [ 0.26076848,  0.23892179,  0.23664738],
       [ 0.29198961,  0.26234032,  0.16505091],
       [ 0.34396535,  0.24206137,  0.16240774]]), indices=array([[3, 0, 5],
       [0, 1, 4],
       [0, 5, 1],
       [1, 3, 5],
       [1, 4, 3]], dtype=int32))
```

Looking just at the first row we get [0.34763842, 0.24879643, 0.12789202], you can confirm these are the 3 largest probabilities in a. You'll also notice [3, 0, 5] are the corresponding indices.

In [ ]:

```
### Print out the top five softmax probabilities for the predictions on the German
### Feel free to use as many code cells as needed.
```

In [129]:

```
def evaluate_non_badge_softmax(X_data, y_data):

    with tf.Session() as sess:
        # Restore variables from disk.
        saver.restore(sess, "./lenet")
        print("Model restored.")

        softmax_logits = tf.nn.softmax(logits)
        top_k = tf.nn.top_k(softmax_logits, k=5)

        #In case it doesnt work uncomment following line
        # my_softmax_logits = sess.run(softmax_logits, feed_dict={x: X_data, y: y_d
        my_top_k = sess.run(top_k, feed_dict={x: X_data, keep_prob: 1.0})

        return my_top_k
```

In [163]:

```python
%matplotlib inline
my_top_k = evaluate_non_badge_softmax(X_test_custom_prep, y_test_custom)


# one time
fig = plt.figure(figsize=(15.0, 8.0), dpi=180)
subplot_grid_n_col = 6 # input + 5 * highest probabilities
subplot_grid_n_row = 5 # 5 sample images
#fig, axs = plt.subplots(len(X_test_custom_prep),4, figsize=(12, 14))
fig.subplots_adjust(hspace = 1.2, wspace=0)
#axs = axs.ravel()

for i, image in enumerate(X_test_custom_prep):
    input_image = X_test_custom[i]
    ax_input = fig.add_subplot(subplot_grid_n_row,subplot_grid_n_col,i * subplot_gr
    ax_input.imshow(input_image.squeeze(), cmap="gray")
    ax_input.set_title("Input {}".format(i))

    # plot top 5 predictions
    for j, (top_pred_label, top_pred_proba) in enumerate(zip(my_top_k[1][i], my_top
        pred_image = (X_train_aug[y_train_aug == top_pred_label][1])

        ax_pred = fig.add_subplot(subplot_grid_n_row,subplot_grid_n_col, (i * subpl
        ax_pred.imshow(pred_image.squeeze(), cmap="gray")
        ax_pred.set_title("Guess {0}\nprobability {1:.1f}%".format(j+1, top_pred_pr

plt.show()
```

INFO:tensorflow:Restoring parameters from ./lenet
Model restored.

Project Writeup

Once you have completed the code implementation, document your results in a project writeup using this template (https://github.com/udacity/CarND-Traffic-Sign-Classifier-Project/blob/master/writeup_template.md) as a guide. The writeup can be in a markdown or pdf file.

> **Note**: Once you have completed all of the code implementations and successfully answered each question above, you may finalize your work by exporting the iPython Notebook as an HTML document. You can do this by using the menu above and navigating to \n", "**File -> Download as -> HTML (.html)**. Include the finished document along with this notebook as your submission.

Step 4 (Optional): Visualize the Neural Network's State with Test Images

This Section is not required to complete but acts as an additional excersise for understaning the output of a neural network's weights. While neural networks can be a great learning device they are often referred to as a black box. We can understand what the weights of a neural network look like better by plotting their feature maps. After successfully training your neural network you can see what it's feature maps look like by plotting the output of the network's weight layers in response to a test stimuli image. From these plotted feature maps, it's possible to see what characteristics of an image the network finds interesting. For a sign, maybe the inner network feature maps react with high activation to the sign's boundary outline or to the contrast in the sign's painted symbol.

Provided for you below is the function code that allows you to get the visualization output of any tensorflow weight layer you want. The inputs to the function should be a stimuli image, one used during training or a new one you provided, and then the tensorflow variable name that represents the layer's state during the training process, for instance if you wanted to see what the LeNet lab's (https://classroom.udacity.com/nanodegrees/nd013/parts/fbf77062-5703-404e-b60c-95b78b2f3f9e/modules/6df7ae49-c61c-4bb2-a23e-6527e69209ec/lessons/601ae704-1035-4287-8b11-e2c2716217ad/concepts/d4aca031-508f-4e0b-b493-e7b706120f81) feature maps looked like for it's second convolutional layer you could enter conv2 as the tf_activation variable.

For an example of what feature map outputs look like, check out NVIDIA's results in their paper End-to-End Deep Learning for Self-Driving Cars (https://devblogs.nvidia.com/parallelforall/deep-learning-self-driving-cars/) in the section Visualization of internal CNN State. NVIDIA was able to show that their network's inner weights had high activations to road boundary lines by comparing feature maps from an image with a clear path to one without. Try experimenting with a similar test to show that your trained network's weights are looking for interesting features, whether it's looking at differences in feature maps from images with or without a sign, or even what feature maps look like in a trained network vs a completely untrained one on the same sign image.

Your output should look something like this (above)

In [ ]:

```
### Visualize your network's feature maps here.
### Feel free to use as many code cells as needed.

# image_input: the test image being fed into the network to produce the feature map
# tf_activation: should be a tf variable name used during your training procedure t
# activation_min/max: can be used to view the activation contrast in more detail, b
# plt_num: used to plot out multiple different weight feature map sets on the same

def outputFeatureMap(image_input, tf_activation, activation_min=-1, activation_max=
    # Here make sure to preprocess your image_input in a way your network expects
    # with size, normalization, ect if needed
    # image_input =
    # Note: x should be the same name as your network's tensorflow data placeholder
    # If you get an error tf_activation is not defined it may be having trouble acc
    activation = tf_activation.eval(session=sess,feed_dict={x : image_input})
    featuremaps = activation.shape[3]
    plt.figure(plt_num, figsize=(15,15))
    for featuremap in range(featuremaps):
        plt.subplot(6,8, featuremap+1) # sets the number of feature maps to show on
        plt.title('FeatureMap ' + str(featuremap)) # displays the feature map numbe
        if activation_min != -1 & activation_max != -1:
            plt.imshow(activation[0,:,:, featuremap], interpolation="nearest", vmin
        elif activation_max != -1:
            plt.imshow(activation[0,:,:, featuremap], interpolation="nearest", vmax
        elif activation_min !=-1:
            plt.imshow(activation[0,:,:, featuremap], interpolation="nearest", vmin
        else:
            plt.imshow(activation[0,:,:, featuremap], interpolation="nearest", cmap
```

What to do to improve, TODO

- experiment different network architectures
- changes dimensions of LeNet
- Tune Hyperparameters
- Augment data by rotate, color, etc.

In [ ]: